From permitted human answers to 100 simulated respondents

100 real respondents

92 evidence questions each
9,198 observed answers

→

100 learned personas

8 numbers per person
Shared model stays fixed

→

8 survey questions

3 reserved accuracy checks
5 new allocation scenarios

→

Predicted responses

5 probabilities per question
500 new single-select draws

Frozen Qwen embeddings → existing 64-feature PCA → shared scorer + each person's eight-number vector

Predictions were saved before final-test answers were opened. Persona vectors are learned coordinates, not invented biographies.